

# Nassim Rafiefard

Code:27

Project: heartdata

## EDA

### Import all libraries

```
In [122...] import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats#
from sklearn.metrics import accuracy_score, recall_score, precision_score, f1_score, confusion_matrix, classification_report
from sklearn.model_selection import cross_val_predict, cross_val_score
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import KFold

import warnings
warnings.filterwarnings("ignore")
```

### Import HeartData

```
In [123...] heartdata=pd.read_csv('/Users/nasimrafie/Documents/data science/Tehran Data/Python/1/Section 1/Datasets/Heart d
```

```
In [123...] heartdata.head()
```

```
Out[123...]
   Age  sex  chest  blood  cholestor  blood  electrocardiographic  heart  exercise  depression  slope  ca  thal  c
   (age  in  pain  pressure  al  sugar  graph  rate  induced  depression  slope  ca  thal  c
   in  year)
0      63    1     1    145.0    233.0     1.0             2.0  150.0     0.0         2.3    3.0  0.0  6.0  0
1      37    1     3    130.0    250.0     0.0             0.0  187.0     0.0         3.5    3.0  0.0  3.0  0
2      41    0     2    130.0    204.0     0.0             2.0  172.0     0.0         1.4    1.0  0.0  3.0  0
3      56    1     2    120.0    236.0     0.0             0.0  178.0     0.0         0.8    1.0  0.0  3.0  0
4      57    0     4    120.0    354.0     0.0             0.0  163.0     1.0         0.6    1.0  0.0  3.0  0
```

```
In [123...] heartdata.tail()
```

```
Out[123...]
   Age  sex  chest  blood  cholestor  blood  electrocardiographic  heart  exercise  depression  slope  ca  thal  c
   (age  in  pain  pressure  al  sugar  graph  rate  induced  depression  slope  ca  thal  c
   in  year)
592    52    1     4    140.0    266.0     0.0             0.0  134.0     1.0         2.0    2.0  NaN  NaN  1
593    43    1     4    140.0    288.0     0.0             0.0  135.0     1.0         2.0    2.0  NaN  NaN  1
594    41    1     4    120.0    336.0     0.0             0.0  118.0     1.0         3.0    2.0  NaN  NaN  1
595    44    1     4    135.0    491.0     0.0             0.0  135.0     0.0         0.0    NaN  NaN  NaN  1
596    49    1     4    150.0    222.0     0.0             0.0  122.0     0.0         2.0    2.0  NaN  NaN  1
```

Data has 587 rows and 14 columns

```
In [123...] heartdata.shape
```

```
Out[123...] (597, 14)
```

```
In [124...] print('There are', heartdata.shape[0], 'rows and', heartdata.shape[1], 'columns in heartdata.')
```

There are 597 rows and 14 columns in heartdata.

```
In [124...] heartdata.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 597 entries, 0 to 596
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Age (age in year)      597 non-null    int64
1   sex                    597 non-null    int64
2   chest pain             597 non-null    int64
3   blood pressure         596 non-null    float64
4   cholestoral            574 non-null    float64
5   blood sugar            589 non-null    float64
6   electrocardiographic  596 non-null    float64
7   heart rate             596 non-null    float64
8   exercise induced       596 non-null    float64
9   depression             597 non-null    float64
10  slope                  407 non-null    float64
11  ca                     303 non-null    float64
12  thal                   329 non-null    float64
13  c                      597 non-null    int64
dtypes: float64(10), int64(4)
memory usage: 65.4 KB
```

```
In [124.. heartdata.describe().T
```

```
Out[124..
```

	count	mean	std	min	25%	50%	75%	max
<b>Age (age in year)</b>	597.0	51.182580	9.074366	28.0	44.0	52.0	58.00	77.0
<b>sex</b>	597.0	0.701843	0.457833	0.0	0.0	1.0	1.00	1.0
<b>chest pain</b>	597.0	3.072027	0.965776	1.0	2.0	3.0	4.00	4.0
<b>blood pressure</b>	596.0	132.129195	17.603812	92.0	120.0	130.0	140.00	200.0
<b>cholestoral</b>	574.0	248.655052	59.784805	85.0	211.0	242.5	278.75	603.0
<b>blood sugar</b>	589.0	0.110357	0.313600	0.0	0.0	0.0	0.00	1.0
<b>electrocardiographic</b>	596.0	0.610738	0.869358	0.0	0.0	0.0	2.00	2.0
<b>heart rate</b>	596.0	144.456376	23.794282	71.0	128.0	146.0	162.00	202.0
<b>exercise induced</b>	596.0	0.315436	0.465080	0.0	0.0	0.0	1.00	1.0
<b>depression</b>	597.0	0.816248	1.067938	0.0	0.0	0.2	1.50	6.2
<b>slope</b>	407.0	1.675676	0.572758	1.0	1.0	2.0	2.00	3.0
<b>ca</b>	303.0	0.693069	1.049212	0.0	0.0	0.0	1.00	9.0
<b>thal</b>	329.0	4.811550	1.928854	3.0	3.0	3.0	7.00	7.0
<b>c</b>	597.0	0.410385	0.492316	0.0	0.0	0.0	1.00	1.0

## Dataset Attributes

Age : age of the patient [years]

Sex : sex of the patient [M: Male, F: Female]

ChestPainType : chest pain type [TA: Typical Angina, ATA: Atypical Angina, NAP: Non-Anginal Pain, ASY: Asymptomatic]

RestingBP : resting blood pressure [mm Hg]

Cholesterol : serum cholesterol [mm/dl]

FastingBS : fasting blood sugar [1: if FastingBS > 120 mg/dl, 0: otherwise]

RestingECG : resting electrocardiogram results [Normal: Normal, ST: having ST-T wave abnormality (T wave inversions and/or ST elevation or depression of > 0.05 mV), LVH: showing probable or definite left ventricular hypertrophy by Estes' criteria]

MaxHR : maximum heart rate achieved [Numeric value between 60 and 202]

ExerciseAngina : exercise-induced angina [Y: Yes, N: No]

Oldpeak : oldpeak = ST [Numeric value measured in depression]

ST\_Slope : the slope of the peak exercise ST segment [Up: upsloping, Flat: flat, Down: downsloping]

HeartDisease : output class [1: heart disease, 0: Normal]

```
In [124.. heartdata.columns
```

```
Out[124...] Index(['Age (age in year)', 'sex', 'chest pain', 'blood pressure',
      'cholesterol ', 'blood sugar', 'electrocardiographic ', 'heart rate',
      'exercise induced', 'depression ', 'slope', 'ca', 'thal', 'c'],
      dtype='object')
```

## Category Counts in Categorical Columns

```
In [125...] cat_col=['sex', 'chest pain', 'electrocardiographic ', 'exercise induced' , 'slope', 'ca', 'thal', 'blood sugar']
for column in cat_col:
    print(heartdata[column].value_counts( dropna=False))
    print('-'*50)
```

```
sex
1    419
0    178
Name: count, dtype: int64
-----
chest pain
4    267
2    156
3    140
1     34
Name: count, dtype: int64
-----
electrocardiographic
0.0    386
2.0    154
1.0     56
NaN      1
Name: count, dtype: int64
-----
exercise induced
0.0    408
1.0    188
NaN      1
Name: count, dtype: int64
-----
slope
2.0    231
NaN    190
1.0    154
3.0     22
Name: count, dtype: int64
-----
ca
NaN    294
0.0    179
1.0     65
2.0     38
3.0     20
9.0      1
Name: count, dtype: int64
-----
thal
NaN    268
3.0    173
7.0    128
6.0     28
Name: count, dtype: int64
-----
blood sugar
0.0    524
1.0     65
NaN      8
Name: count, dtype: int64
-----
```

```
In [125...] for column in cat_col:
    print(heartdata[column].value_counts(normalize=True, dropna=False))
    print('-'*50)
```

```

sex
1    0.701843
0    0.298157
Name: proportion, dtype: float64
-----
chest pain
4    0.447236
2    0.261307
3    0.234506
1    0.056951
Name: proportion, dtype: float64
-----
electrocardiographic
0.0    0.646566
2.0    0.257956
1.0    0.093802
NaN    0.001675
Name: proportion, dtype: float64
-----
exercise induced
0.0    0.683417
1.0    0.314908
NaN    0.001675
Name: proportion, dtype: float64
-----
slope
2.0    0.386935
NaN    0.318258
1.0    0.257956
3.0    0.036851
Name: proportion, dtype: float64
-----
ca
NaN    0.492462
0.0    0.299832
1.0    0.108878
2.0    0.063652
3.0    0.033501
9.0    0.001675
Name: proportion, dtype: float64
-----
thal
NaN    0.448911
3.0    0.289782
7.0    0.214405
6.0    0.046901
Name: proportion, dtype: float64
-----
blood sugar
0.0    0.877722
1.0    0.108878
NaN    0.013400
Name: proportion, dtype: float64
-----

```

**Observation** In 'ca' column value 9.0 is observed which should be 0.0 and has been typed 9.0 by mistake.

```
In [125...] heartdata['ca'][heartdata['ca']==9]=0
```

```
In [125...] heartdata['ca'].value_counts(normalize=True, dropna=False)
```

```
Out[125...] ca
NaN    0.492462
0.0    0.301508
1.0    0.108878
2.0    0.063652
3.0    0.033501
Name: proportion, dtype: float64
```

## Missing Values

```
In [126...] heartdata.isnull().sum()
```

```
Out[126...] Age (age in year)      0
sex          0
chest pain   0
blood pressure 1
cholestorl   23
blood sugar   8
electrocardiographic 1
heart rate    1
exercise induced 1
depression    0
slope        190
ca           294
thal         268
c            0
dtype: int64
```

## Univariate Observation

```
In [126...] def histogram_boxplot(feature, figsize=(15, 10), bins="auto"):
    """ Boxplot and histogram combined
    feature: 1-d feature array
    figsize: size of fig (default (15, 10))
    bins: number of bins (default "auto")
    """
    f, (ax_box, ax_hist) = plt.subplots(
        nrows=2, # Number of rows of the subplot grid
        sharex=True, # The X-axis will be shared among all the subplots
        gridspec_kw={"height_ratios": (.25, .75)},
        figsize=figsize
    )

    # Creating the subplots
    # Boxplot will be created and the mean value of the column will be indicated using some symbol
    sns.boxplot(x=feature, ax=ax_box, showmeans=True, color='red')





    # For histogram
    sns.histplot(x=feature, kde=False, ax=ax_hist, bins=bins)
    ax_hist.axvline(np.mean(feature), color='g', linestyle='--') # Add mean to the histogram
    ax_hist.axvline(np.median(feature), color='black', linestyle='-') # Add median to the histogram





    plt.show()
```


```
In [126...] num_coll = heartdata.drop(columns=cat_col)
num_col=num_coll.drop('c',axis=1)
num_col.columns
```

```
Out[126...] Index(['Age (age in year)', 'blood pressure', 'cholestorl ', 'heart rate',
      'depression '],
      dtype='object')
```

```
In [126...] for i in num_col:
    p = histogram_boxplot(heartdata[i])
    plt.show()
```







## Observation

In [ ]:

```
In [127...] def bar_perc(data, z):
    total = len(data[z]) # Length of the column
    plt.figure(figsize = (15, 5))


    # Convert the column to a categorical data type
    data[z] = data[z].astype('category')

    ax = sns.countplot(x=z, data=data, hue=z, palette='Paired', order=data[z].value_counts().index)


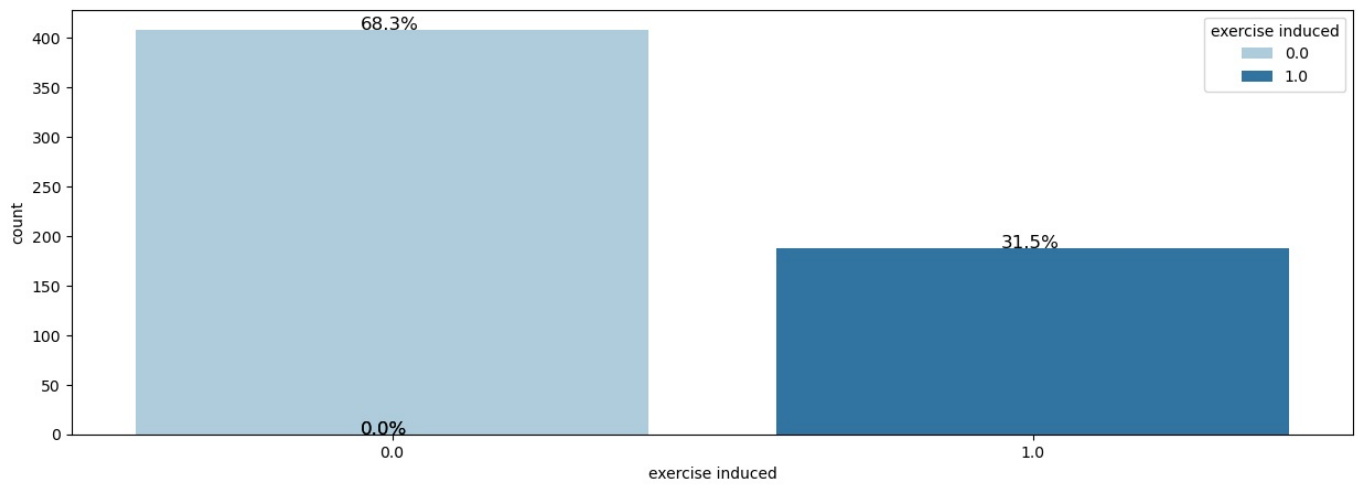

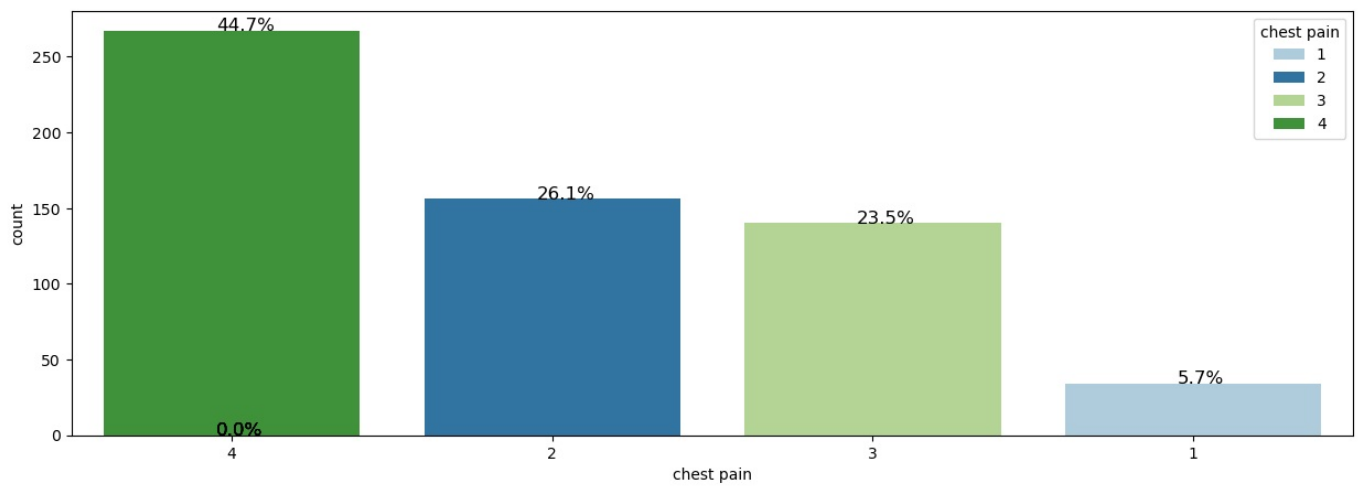
    for p in ax.patches:
        percentage = '{:.1f}%'.format(100 * p.get_height() / total) # Percentage of each class
        x = p.get_x() + p.get_width() / 2 - 0.05 # Width of the plot
        y = p.get_y() + p.get_height() # Height of the plot
        ax.annotate(percentage, (x, y), size = 12) # Annotate the percentage

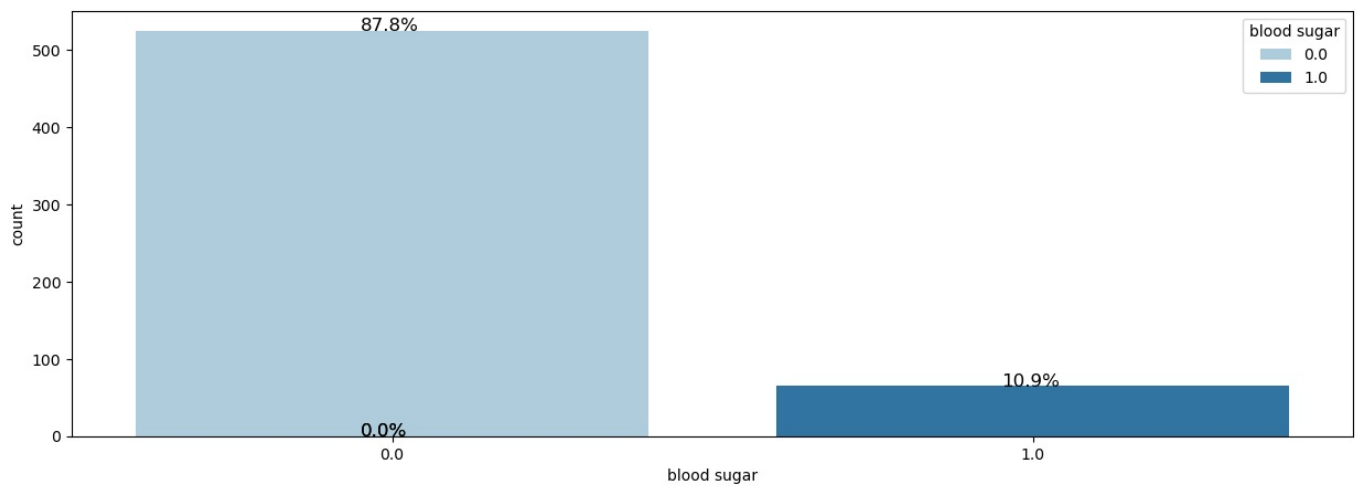


    plt.show() # Display the plot
```

```
In [127...] for i in cat_col:
    p = bar_perc(heartdata,i)
    plt.show()
```









## Multivariate Observation

In [127... heartdata.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 597 entries, 0 to 596
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Age (age in year)      597 non-null    int64
1   sex                    597 non-null    category
2   chest pain             597 non-null    category
3   blood pressure         596 non-null    float64
4   cholestoral            574 non-null    float64
5   blood sugar            589 non-null    category
6   electrocardiographic  596 non-null    category
7   heart rate             596 non-null    float64
8   exercise induced       596 non-null    category
9   depression             597 non-null    float64
10  slope                  407 non-null    category
11  ca                     303 non-null    category
12  thal                   329 non-null    category
13  c                      597 non-null    int64
dtypes: category(8), float64(4), int64(2)
memory usage: 33.9 KB
```

```
In [127.. corr=heartdata.corr()
corr
```


Out[127..

	Age (age in year)	sex	chest pain	blood pressure	cholestoral	blood sugar	electrocardiographic	heart rate	exercise induced	c
Age (age in year)	1.000000	-0.062397	0.147064	0.238490	0.123624	0.176286	0.260132	-0.303596	0.155862	
sex	-0.062397	1.000000	0.120748	0.010620	-0.076399	0.038030	-0.034982	-0.088691	0.148814	
chest pain	0.147064	0.120748	1.000000	0.021586	0.104111	0.001428	0.073320	-0.322748	0.438328	
blood pressure	0.238490	0.010620	0.021586	1.000000	0.105189	0.136097	0.077768	-0.117829	0.136658	
cholestoral	0.123624	-0.076399	0.104111	0.105189	1.000000	0.054867	0.088498	-0.076064	0.117111	
blood sugar	0.176286	0.038030	0.001428	0.136097	0.054867	1.000000	0.111847	-0.005236	0.063431	
electrocardiographic	0.260132	-0.034982	0.073320	0.077768	0.088498	0.111847	1.000000	0.052515	0.071970	
heart rate	-0.303596	-0.088691	-0.322748	-0.117829	-0.076064	-0.005236	0.052515	1.000000	-0.374642	
exercise induced	0.155862	0.148814	0.438328	0.136658	0.117111	0.063431	0.071970	-0.374642	1.000000	
depression	0.253305	0.095716	0.277695	0.185216	0.065998	0.050842	0.175329	-0.259880	0.426849	
slope	0.078979	0.075835	0.209141	0.126015	0.047846	0.058897	0.032245	-0.402652	0.332025	
ca	0.364036	0.090833	0.227668	0.093548	0.123661	0.148741	0.136486	-0.253548	0.140423	
thal	0.105296	0.349134	0.245214	0.133696	0.011964	0.069128	-0.012682	-0.302562	0.320352	
c	0.216430	0.268343	0.463527	0.142178	0.145802	0.090071	0.137410	-0.342209	0.504280	

```
In [128.. ## Check for correlation among numerical variables
#corr = heartdata.corr()

# Plot the mapp
plt.figure(figsize = (14, 10))
sns.heatmap(corr, annot = True, cmap = 'coolwarm', fmt = ".1f")
```

Out[128.. <Axes: >



## Observation

Chest pain, exercise-induced angina, depression, calcium levels (ca), and thalassemia (thal) all appear to have a correlation with c, with values around 0.5.

## Missing value

```
In [128]: heartdata.isnull().sum()
```

```
Out[128.. Age (age in year)          0
sex          0
chest pain   0
blood pressure 1
cholestorl   23
blood sugar   8
electrocardiographic 1
heart rate    1
exercise induced 1
depression    0
slope        190
ca           294
thal         268
c            0
dtype: int64
```

## Outlier Values

### Number of outliers based on IQR and Z-score

```
In [129.. for variable in num_col:
print(variable)
Q1=heartdata[variable].quantile(0.25)
Q3=heartdata[variable].quantile(0.75)
IQR=Q3-Q1
zscore=np.abs(stats.zscore(heartdata[variable]))

# print ('min=',heartdata[variable].min(),'max=',heartdata[variable].max(),'median=',heartdata[variable].med.
print('number of outliers with IQR=',heartdata[(heartdata[variable]<Q1-1.5*IQR)|(heartdata[variable]>Q3+1.5
print('number of outliers with zscore=',heartdata[zscore>3].shape[0])
print( zscore[zscore>3])
```

```
Age (age in year)
number of outliers with IQR= 0
number of outliers with zscore= 0
Series([], Name: Age (age in year), dtype: float64)
blood pressure
number of outliers with IQR= 17
number of outliers with zscore= 0
Series([], Name: blood pressure, dtype: float64)
cholestorl
number of outliers with IQR= 19
number of outliers with zscore= 0
Series([], Name: cholestoral , dtype: float64)
heart rate
number of outliers with IQR= 1
number of outliers with zscore= 0
Series([], Name: heart rate, dtype: float64)
depression
number of outliers with IQR= 11
number of outliers with zscore= 6
101    3.171150
514    5.045489
520    4.483187
527    3.171150
581    3.358583
591    3.920885
Name: depression , dtype: float64
```

Fill outliers which are found by Z-score with null values and the rest with upper and lower amount of data.

```
In [129.. def treat_outliers(df, col):
# Calculate IQR and Z-score
Q1 = df[col].quantile(0.25)
Q3 = df[col].quantile(0.75)
IQR = Q3 - Q1
lower_whisker = Q1 - 1.5 * IQR
upper_whisker = Q3 + 1.5 * IQR
zscore = np.abs(stats.zscore(df[col]))

# Loop through each value and apply conditions
for idx, value in enumerate(df[col]):
```

```

if (value < lower_whisker or value > upper_whisker) and zscore[idx] > 3:
    # Set to NaN if outlier in both IQR and Z-score
    df.at[idx, col] = np.nan
elif value < lower_whisker:
    # Clip to lower limit if only an IQR outlier
    df.at[idx, col] = lower_whisker
elif value > upper_whisker:
    # Clip to upper limit if only an IQR outlier
    df.at[idx, col] = upper_whisker

return df

# Apply the function to each relevant column
num_cols = ['blood pressure', 'cholesterol', 'heart rate', 'depression']
for col in num_cols:
    heartdata = treat_outliers(heartdata, col)

```

Outlier found based on Z-score are added to null values.

In [129.. heartdata.isnull().sum()

```

Out[129.. Age (age in year)      0
sex                          0
chest pain                   0
blood pressure                1
cholesterol                   23
blood sugar                   8
electrocardiographic         1
heart rate                    1
exercise induced              1
depression                    6
slope                         190
ca                            294
thal                          268
c                              0
dtype: int64


```





No more outliers are detected in box-plot





```

In [130.. for i in num_col:
    p = histogram_boxplot(heartdata[i])
    plt.show()

```







Missing



```
In [130]: heartdata.isnull().sum()

Out[130]: Age (age in year)      0
sex                             0
chest pain                      0
blood pressure                  1
cholesterol                     23
blood sugar                     8
electrocardiographic           1
heart rate                      1
exercise induced                1
depression                      6
slope                           190
ca                              294
thal                           268
c                               0
dtype: int64
```

```
In [130]: pd.DataFrame({'Count':heartdata.isnull().sum()[heartdata.isnull().sum()>0], 'Percentage':heartdata.isnull().sum
```

Out[130]:

	Count	Percentage
blood pressure	1	0.167504
cholesterol	23	3.852596
blood sugar	8	1.340034
electrocardiographic	1	0.167504
heart rate	1	0.167504
exercise induced	1	0.167504
depression	6	1.005025
slope	190	31.825796
ca	294	49.246231
thal	268	44.891122

Observation

ca and thal columns have the highest percentage (almost 50%) of missing values.

blood pressure, electrocardiographic, heart rate and exercise induced have the lowest percentage of missing values.

There is no column with more than 50% of missing values.

Finding rows with more than 4 null features to drop

```
In [131]: heartdata.columns
```

```
Out[131]: Index(['Age (age in year)', 'sex', 'chest pain', 'blood pressure',
               'cholesterol ', 'blood sugar', 'electrocardiographic ', 'heart rate',
               'exercise induced', 'depression ', 'slope', 'ca', 'thal', 'c'],
              dtype='object')
```

```
In [131]: heartdata.loc[heartdata['blood pressure'].isnull()==True]
```

Out[131]:

	Age (age in year)	sex	chest pain	blood pressure	cholesterol	blood sugar	electrocardiographic	heart rate	exercise induced	depression	slope	ca	thal	c
347	48	0	2	NaN	308.0	0.0	1.0	NaN	NaN	2.0	1.0	NaN	NaN	0

```
In [131]: heartdata.loc[heartdata['electrocardiographic '].isnull()==True]
```

Out[131]:

	Age (age in year)	sex	chest pain	blood pressure	cholesterol	blood sugar	electrocardiographic	heart rate	exercise induced	depression	slope	ca	thal	c
562	55	1	1	140.0	295.0	0.0	NaN	136.0	0.0	0.0	NaN	NaN	NaN	1

```
In [131]: heartdata.loc[heartdata['heart rate'].isnull()==True]
```

Out[131..

	Age (age in year)	sex	chest pain	blood pressure	cholestorol	blood sugar	electrocardiographic	heart rate	exercise induced	depression	slope	ca	thal	c	
347	48	0	2	NaN	308.0	0.0		1.0	NaN	NaN	2.0	1.0	NaN	NaN	0

In [131..

```
heartdata.loc[heartdata['exercise induced'].isnull()==True]
```

Out[131..

	Age (age in year)	sex	chest pain	blood pressure	cholestorol	blood sugar	electrocardiographic	heart rate	exercise induced	depression	slope	ca	thal	c	
347	48	0	2	NaN	308.0	0.0		1.0	NaN	NaN	2.0	1.0	NaN	NaN	0

In [132..

```
heartdata.loc[heartdata['blood sugar'].isnull()==True]
```

Out[132..

	Age (age in year)	sex	chest pain	blood pressure	cholestorol	blood sugar	electrocardiographic	heart rate	exercise induced	depression	slope	ca	thal	c
183	53	0	2	113.0	380.375	NaN	0.0	127.0	0.0	0.0	NaN	NaN	NaN	0
213	46	1	3	150.0	163.000	NaN	0.0	116.0	0.0	0.0	NaN	NaN	NaN	0
289	49	1	4	120.0	297.000	NaN	0.0	132.0	0.0	1.0	2.0	NaN	NaN	0
301	56	0	3	130.0	219.000	NaN	1.0	164.0	0.0	0.0	NaN	NaN	7.0	0
314	38	0	2	120.0	275.000	NaN	0.0	129.0	0.0	0.0	NaN	NaN	NaN	0
316	54	0	2	140.0	309.000	NaN	1.0	140.0	0.0	0.0	NaN	NaN	NaN	0
421	40	1	4	120.0	380.375	NaN	0.0	152.0	1.0	1.0	2.0	NaN	6.0	1
431	41	1	4	120.0	237.000	NaN	0.0	138.0	1.0	1.0	2.0	NaN	NaN	1

In [132..

```
heartdata.loc[heartdata['cholestorol '].isnull()==True].head()
```

Out[132..

	Age (age in year)	sex	chest pain	blood pressure	cholestorol	blood sugar	electrocardiographic	heart rate	exercise induced	depression	slope	ca	thal	c
166	54	1	3	150.0	NaN	0.0	0.0	122.0	0.0	0.0	NaN	NaN	NaN	0
185	56	1	3	130.0	NaN	0.0	0.0	114.0	0.0	0.0	NaN	NaN	NaN	0
189	48	0	2	120.0	NaN	1.0	1.0	148.0	0.0	0.0	NaN	NaN	NaN	0
204	43	0	3	150.0	NaN	0.0	0.0	175.0	0.0	0.0	NaN	NaN	3.0	0
221	49	0	2	110.0	NaN	0.0	0.0	160.0	0.0	0.0	NaN	NaN	NaN	0

Blood pressure, heart rate and exercise induced's only missing value are from the same index No. 347, where Thal and Ca is also null. That is why it makes sense to delete the data. There is one line with 5 missing values.

In [132..

```
heartdata = heartdata.drop(347).reset_index(drop=True)
```

line 347 removed

In [132..

```
pd.DataFrame({'Count':heartdata.isnull().sum()[heartdata.isnull().sum()>0], 'Percentage':heartdata.isnull().sum
```

Out[132..

	Count	Percentage
cholestorol	23	3.859060
blood sugar	8	1.342282
electrocardiographic	1	0.167785
depression	6	1.006711
slope	190	31.879195
ca	293	49.161074
thal	267	44.798658

Exercise doesnt have null value any more.

Check if null values from thal,ca and slope are from same index:

```
In [133...] heartdata.loc[heartdata['thal'].isnull()==True,'ca'].value_counts(normalize=True,dropna=False)
```

```
Out[133...] ca
NaN      0.981273
0.0      0.018727
1.0      0.000000
2.0      0.000000
3.0      0.000000
Name: proportion, dtype: float64
```

```
In [133...] heartdata.loc[heartdata['slope'].isnull()==True,'ca'].value_counts(normalize=True, dropna=False)
```

```
Out[133...] ca
NaN      0.989474
0.0      0.010526
1.0      0.000000
2.0      0.000000
3.0      0.000000
Name: proportion, dtype: float64
```

```
In [133...] heartdata.loc[heartdata['slope'].isnull()==True,'thal'].value_counts(normalize=True,dropna=False)
```

```
Out[133...] thal
NaN      0.910526
7.0      0.036842
6.0      0.031579
3.0      0.021053
Name: proportion, dtype: float64
```

```
In [133...] heartdata.loc[heartdata['cholesterol '].isnull()==True,'ca'].value_counts(normalize=True, dropna=False)
```

```
Out[133...] ca
NaN      0.956522
0.0      0.043478
1.0      0.000000
2.0      0.000000
3.0      0.000000
Name: proportion, dtype: float64
```

There seems to be a strong pattern in missing values, as wherever the ca column has missing data the thal, cholesterol and slope columns also have missing values.

```
In [134...] nullcholestral=heartdata.loc[heartdata['cholesterol '].isnull()==True]
nullcholestral
```

Out[134..

	Age (age in year)	sex	chest pain	blood pressure	cholesterol	blood sugar	electrocardiographic	heart rate	exercise induced	depression	slope	ca	thal	c
166	54	1	3	150.0	NaN	0.0	0.0	122.0	0.0	0.0	NaN	NaN	NaN	0
185	56	1	3	130.0	NaN	0.0	0.0	114.0	0.0	0.0	NaN	NaN	NaN	0
189	48	0	2	120.0	NaN	1.0	1.0	148.0	0.0	0.0	NaN	NaN	NaN	0
204	43	0	3	150.0	NaN	0.0	0.0	175.0	0.0	0.0	NaN	NaN	3.0	0
221	49	0	2	110.0	NaN	0.0	0.0	160.0	0.0	0.0	NaN	NaN	NaN	0
251	39	1	2	120.0	NaN	0.0	1.0	146.0	0.0	2.0	1.0	NaN	NaN	0
257	39	1	2	130.0	NaN	0.0	0.0	120.0	0.0	0.0	NaN	NaN	NaN	0
258	48	1	2	100.0	NaN	0.0	0.0	100.0	0.0	0.0	NaN	NaN	NaN	0
262	49	1	4	140.0	NaN	0.0	0.0	130.0	0.0	0.0	NaN	NaN	NaN	0
265	49	0	2	110.0	NaN	0.0	0.0	160.0	0.0	0.0	NaN	NaN	NaN	0
267	52	0	2	140.0	NaN	0.0	0.0	140.0	0.0	0.0	NaN	NaN	NaN	0
272	40	1	3	140.0	NaN	0.0	0.0	188.0	0.0	0.0	NaN	NaN	NaN	0
286	45	1	3	135.0	NaN	0.0	0.0	110.0	0.0	0.0	NaN	NaN	NaN	0
305	29	1	2	140.0	NaN	0.0	0.0	170.0	0.0	0.0	NaN	NaN	NaN	0
332	47	0	3	130.0	NaN	0.0	0.0	145.0	0.0	2.0	2.0	NaN	NaN	0
336	45	0	2	170.0	NaN	0.0	0.0	180.0	0.0	0.0	NaN	NaN	NaN	0
345	59	1	4	140.0	NaN	0.0	0.0	140.0	0.0	0.0	NaN	0.0	NaN	0
346	53	1	2	120.0	NaN	0.0	0.0	132.0	0.0	0.0	NaN	NaN	NaN	0
489	54	1	4	140.0	NaN	0.0	0.0	118.0	1.0	0.0	NaN	NaN	NaN	1
491	38	1	4	110.0	NaN	0.0	0.0	150.0	1.0	1.0	2.0	NaN	NaN	1
504	52	1	4	170.0	NaN	0.0	0.0	126.0	1.0	1.5	2.0	NaN	NaN	1
549	66	1	4	140.0	NaN	0.0	0.0	94.0	1.0	1.0	2.0	NaN	NaN	1
565	59	1	4	130.0	NaN	0.0	0.0	125.0	0.0	0.0	NaN	NaN	NaN	1

In [134..

```
nullcholestralslope=nullcholestral.loc[nullcholestral['slope'].isnull()==True]  
nullcholestralslope
```

Out[134..

	Age (age in year)	sex	chest pain	blood pressure	cholesterol	blood sugar	electrocardiographic	heart rate	exercise induced	depression	slope	ca	thal	c
166	54	1	3	150.0	NaN	0.0	0.0	122.0	0.0	0.0	NaN	NaN	NaN	0
185	56	1	3	130.0	NaN	0.0	0.0	114.0	0.0	0.0	NaN	NaN	NaN	0
189	48	0	2	120.0	NaN	1.0	1.0	148.0	0.0	0.0	NaN	NaN	NaN	0
204	43	0	3	150.0	NaN	0.0	0.0	175.0	0.0	0.0	NaN	NaN	3.0	0
221	49	0	2	110.0	NaN	0.0	0.0	160.0	0.0	0.0	NaN	NaN	NaN	0
257	39	1	2	130.0	NaN	0.0	0.0	120.0	0.0	0.0	NaN	NaN	NaN	0
258	48	1	2	100.0	NaN	0.0	0.0	100.0	0.0	0.0	NaN	NaN	NaN	0
262	49	1	4	140.0	NaN	0.0	0.0	130.0	0.0	0.0	NaN	NaN	NaN	0
265	49	0	2	110.0	NaN	0.0	0.0	160.0	0.0	0.0	NaN	NaN	NaN	0
267	52	0	2	140.0	NaN	0.0	0.0	140.0	0.0	0.0	NaN	NaN	NaN	0
272	40	1	3	140.0	NaN	0.0	0.0	188.0	0.0	0.0	NaN	NaN	NaN	0
286	45	1	3	135.0	NaN	0.0	0.0	110.0	0.0	0.0	NaN	NaN	NaN	0
305	29	1	2	140.0	NaN	0.0	0.0	170.0	0.0	0.0	NaN	NaN	NaN	0
336	45	0	2	170.0	NaN	0.0	0.0	180.0	0.0	0.0	NaN	NaN	NaN	0
345	59	1	4	140.0	NaN	0.0	0.0	140.0	0.0	0.0	NaN	0.0	NaN	0
346	53	1	2	120.0	NaN	0.0	0.0	132.0	0.0	0.0	NaN	NaN	NaN	0
489	54	1	4	140.0	NaN	0.0	0.0	118.0	1.0	0.0	NaN	NaN	NaN	1
565	59	1	4	130.0	NaN	0.0	0.0	125.0	0.0	0.0	NaN	NaN	NaN	1

In [134..

```
nullcholestralslopeca=nullcholestralslope.loc[nullcholestralslope['ca'].isnull()==True]
```

```
nullcholestralslopeca
```

Out[134...

	Age (age in year)	sex	chest pain	blood pressure	cholestorl	blood sugar	electrocardiographic	heart rate	exercise induced	depression	slope	ca	thal	c
166	54	1	3	150.0	NaN	0.0	0.0	122.0	0.0	0.0	NaN	NaN	NaN	0
185	56	1	3	130.0	NaN	0.0	0.0	114.0	0.0	0.0	NaN	NaN	NaN	0
189	48	0	2	120.0	NaN	1.0	1.0	148.0	0.0	0.0	NaN	NaN	NaN	0
204	43	0	3	150.0	NaN	0.0	0.0	175.0	0.0	0.0	NaN	NaN	3.0	0
221	49	0	2	110.0	NaN	0.0	0.0	160.0	0.0	0.0	NaN	NaN	NaN	0
257	39	1	2	130.0	NaN	0.0	0.0	120.0	0.0	0.0	NaN	NaN	NaN	0
258	48	1	2	100.0	NaN	0.0	0.0	100.0	0.0	0.0	NaN	NaN	NaN	0
262	49	1	4	140.0	NaN	0.0	0.0	130.0	0.0	0.0	NaN	NaN	NaN	0
265	49	0	2	110.0	NaN	0.0	0.0	160.0	0.0	0.0	NaN	NaN	NaN	0
267	52	0	2	140.0	NaN	0.0	0.0	140.0	0.0	0.0	NaN	NaN	NaN	0
272	40	1	3	140.0	NaN	0.0	0.0	188.0	0.0	0.0	NaN	NaN	NaN	0
286	45	1	3	135.0	NaN	0.0	0.0	110.0	0.0	0.0	NaN	NaN	NaN	0
305	29	1	2	140.0	NaN	0.0	0.0	170.0	0.0	0.0	NaN	NaN	NaN	0
336	45	0	2	170.0	NaN	0.0	0.0	180.0	0.0	0.0	NaN	NaN	NaN	0
346	53	1	2	120.0	NaN	0.0	0.0	132.0	0.0	0.0	NaN	NaN	NaN	0
489	54	1	4	140.0	NaN	0.0	0.0	118.0	1.0	0.0	NaN	NaN	NaN	1
565	59	1	4	130.0	NaN	0.0	0.0	125.0	0.0	0.0	NaN	NaN	NaN	1

In [134...

```
nullcholestralslopeca Thal = nullcholestralslopeca.loc[nullcholestralslopeca['thal'].isnull()==True]  
nullcholestralslopeca Thal
```

Out[134...

	Age (age in year)	sex	chest pain	blood pressure	cholestorl	blood sugar	electrocardiographic	heart rate	exercise induced	depression	slope	ca	thal	c
166	54	1	3	150.0	NaN	0.0	0.0	122.0	0.0	0.0	NaN	NaN	NaN	0
185	56	1	3	130.0	NaN	0.0	0.0	114.0	0.0	0.0	NaN	NaN	NaN	0
189	48	0	2	120.0	NaN	1.0	1.0	148.0	0.0	0.0	NaN	NaN	NaN	0
221	49	0	2	110.0	NaN	0.0	0.0	160.0	0.0	0.0	NaN	NaN	NaN	0
257	39	1	2	130.0	NaN	0.0	0.0	120.0	0.0	0.0	NaN	NaN	NaN	0
258	48	1	2	100.0	NaN	0.0	0.0	100.0	0.0	0.0	NaN	NaN	NaN	0
262	49	1	4	140.0	NaN	0.0	0.0	130.0	0.0	0.0	NaN	NaN	NaN	0
265	49	0	2	110.0	NaN	0.0	0.0	160.0	0.0	0.0	NaN	NaN	NaN	0
267	52	0	2	140.0	NaN	0.0	0.0	140.0	0.0	0.0	NaN	NaN	NaN	0
272	40	1	3	140.0	NaN	0.0	0.0	188.0	0.0	0.0	NaN	NaN	NaN	0
286	45	1	3	135.0	NaN	0.0	0.0	110.0	0.0	0.0	NaN	NaN	NaN	0
305	29	1	2	140.0	NaN	0.0	0.0	170.0	0.0	0.0	NaN	NaN	NaN	0
336	45	0	2	170.0	NaN	0.0	0.0	180.0	0.0	0.0	NaN	NaN	NaN	0
346	53	1	2	120.0	NaN	0.0	0.0	132.0	0.0	0.0	NaN	NaN	NaN	0
489	54	1	4	140.0	NaN	0.0	0.0	118.0	1.0	0.0	NaN	NaN	NaN	1
565	59	1	4	130.0	NaN	0.0	0.0	125.0	0.0	0.0	NaN	NaN	NaN	1

In [134...

```
nullcholestralslopeca Thal.index
```

Out[134...

```
Index([166, 185, 189, 221, 257, 258, 262, 265, 267, 272, 286, 305, 336, 346,  
      489, 565],  
      dtype='int64')
```

In above indexes all four cholestoral, Thal,ca and slope are null. These 16 indexes were chosen to be deleted from data

Totally 17 indexes are dropped.

In [135...

```
heartdata = heartdata.drop(nullcholestralslopeca Thal.index).reset_index(drop=True)
```

```
In [135]: pd.DataFrame({'Count':heartdata.isnull().sum()[heartdata.isnull().sum()>0], 'Percentage':heartdata.isnull().sum
```

Out[135]:

	Count	Percentage
cholestorl	7	1.206897
blood sugar	8	1.379310
electrocardiographic	1	0.172414
depression	6	1.034483
slope	174	30.000000
ca	277	47.758621
thal	251	43.275862

```
In [135]: corr
```

Out[135]:

	Age (age in year)	sex	chest pain	blood pressure	cholestorl	blood sugar	electrocardiographic	heart rate	exercise induced
Age (age in year)	1.000000	-0.062397	0.147064	0.238490	0.123624	0.176286	0.260132	-0.303596	0.155862
sex	-0.062397	1.000000	0.120748	0.010620	-0.076399	0.038030	-0.034982	-0.088691	0.148814
chest pain	0.147064	0.120748	1.000000	0.021586	0.104111	0.001428	0.073320	-0.322748	0.438328
blood pressure	0.238490	0.010620	0.021586	1.000000	0.105189	0.136097	0.077768	-0.117829	0.136658
cholestorl	0.123624	-0.076399	0.104111	0.105189	1.000000	0.054867	0.088498	-0.076064	0.117111
blood sugar	0.176286	0.038030	0.001428	0.136097	0.054867	1.000000	0.111847	-0.005236	0.063431
electrocardiographic	0.260132	-0.034982	0.073320	0.077768	0.088498	0.111847	1.000000	0.052515	0.071970
heart rate	-0.303596	-0.088691	-0.322748	-0.117829	-0.076064	-0.005236	0.052515	1.000000	-0.374642
exercise induced	0.155862	0.148814	0.438328	0.136658	0.117111	0.063431	0.071970	-0.374642	1.000000
depression	0.253305	0.095716	0.277695	0.185216	0.065998	0.050842	0.175329	-0.259880	0.426849
slope	0.078979	0.075835	0.209141	0.126015	0.047846	0.058897	0.032245	-0.402652	0.332025
ca	0.364036	0.090833	0.227668	0.093548	0.123661	0.148741	0.136486	-0.253548	0.140423
thal	0.105296	0.349134	0.245214	0.133696	0.011964	0.069128	-0.012682	-0.302562	0.320352
c	0.216430	0.268343	0.463527	0.142178	0.145802	0.090071	0.137410	-0.342209	0.504280

```
In [135]: heartdata.columns
```

```
Out[135]: Index(['Age (age in year)', 'sex', 'chest pain', 'blood pressure',  
          'cholestorl ', 'blood sugar', 'electrocardiographic ', 'heart rate',  
          'exercise induced', 'depression ', 'slope', 'ca', 'thal', 'c'],  
          dtype='object')
```

```
In [136]: heartdata.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 580 entries, 0 to 579  
Data columns (total 14 columns):  
#   Column              Non-Null Count  Dtype  
---  -  
0   Age (age in year)    580 non-null    int64  
1   sex                  580 non-null    category  
2   chest pain           580 non-null    category  
3   blood pressure       580 non-null    float64  
4   cholestorl           573 non-null    float64  
5   blood sugar          572 non-null    category  
6   electrocardiographic 579 non-null    category  
7   heart rate           580 non-null    float64  
8   exercise induced     580 non-null    category  
9   depression           574 non-null    float64  
10  slope                406 non-null    category  
11  ca                   303 non-null    category  
12  thal                 329 non-null    category  
13  c                    580 non-null    int64  
dtypes: category(8), float64(4), int64(2)  
memory usage: 33.0 KB
```

Determine the mode for categorical data grouped by sex , so that any null values can be appropriately filled with the most frequent value.

Null values in the 'blood sugar' column are filled with the mode of Blood sugar, calculated within groups based on sex.

```
In [136.. heartdata.groupby(['sex'])['blood sugar'].agg(lambda x: x.mode().iloc[0])
```

```
Out[136.. sex
0      0.0
1      0.0
Name: blood sugar, dtype: category
Categories (2, float64): [0.0, 1.0]
```

```
In [136.. heartdata['blood sugar'] = heartdata['blood sugar'].fillna(
    heartdata.groupby(['sex'])['blood sugar'].transform(lambda x: x.mode().iloc[0] ))
```

Null values in the 'slope' column are filled with the mode of slope, calculated within groups based on sex.

```
In [136.. heartdata.groupby(['sex'])['slope'].agg(lambda x: x.mode().iloc[0])
```

```
Out[136.. sex
0      2.0
1      2.0
Name: slope, dtype: category
Categories (3, float64): [1.0, 2.0, 3.0]
```

```
In [137.. heartdata['slope'] = heartdata['slope'].fillna(
    heartdata.groupby(['sex'])['slope'].transform(lambda x: x.mode().iloc[0] ))
```

Null values in the 'ca' column are filled with the mode of Ca, calculated within groups based on sex.

```
In [137.. heartdata.groupby(['sex'])['ca'].agg(lambda x: x.mode().iloc[0])
```

```
Out[137.. sex
0      0.0
1      0.0
Name: ca, dtype: category
Categories (4, float64): [0.0, 1.0, 2.0, 3.0]
```

```
In [137.. heartdata['ca'] = heartdata['ca'].fillna(
    heartdata.groupby(['sex'])['ca'].transform(lambda x: x.mode().iloc[0] ))
```

Null values in the 'thal' column are filled with the mode of thal, calculated within groups based on sex.

```
In [137.. heartdata.groupby(['sex'])['thal'].agg(lambda x: x.mode().iloc[0])
```

```
Out[137.. sex
0      3.0
1      7.0
Name: thal, dtype: category
Categories (3, float64): [3.0, 6.0, 7.0]
```

```
In [138.. heartdata['thal'] = heartdata['thal'].fillna(
    heartdata.groupby(['sex'])['thal'].transform(lambda x: x.mode().iloc[0] ))
```

Null values in the electrocardiographic column are filled with the mode of electrocardiographic, calculated within groups based on sex.

```
In [138.. heartdata.groupby(['sex'])['electrocardiographic '].agg(lambda x: x.mode().iloc[0])
```

```
Out[138.. sex
0      0.0
1      0.0
Name: electrocardiographic , dtype: category
Categories (3, float64): [0.0, 1.0, 2.0]
```

```
In [138.. heartdata['electrocardiographic '] = heartdata['electrocardiographic '].fillna(
    heartdata.groupby(['sex'])['electrocardiographic '].transform(lambda x: x.mode().iloc[0] ))
```

Null values in the Blood Pressure column are filled with the mode of Blood Pressure, calculated within groups based on sex.

```
In [138.. heartdata.groupby(['sex'])['blood pressure'].agg(lambda x: x.mode().iloc[0])
```

```
Out[138.. sex
0     130.0
1     120.0
Name: blood pressure, dtype: float64
```

```
In [139.. heartdata['blood pressure'] = heartdata['blood pressure'].fillna(
    heartdata.groupby(['sex'])['blood pressure'].transform(lambda x: x.mode().iloc[0] ))
```

Null values in the cholestral column are filled with the mean of cholestral, calculated within groups based on

sex.

```
In [139.. heartdata.groupby(['sex'])['cholesterol '].mean()
```

```
Out[139.. sex
0    252.913971
1    243.860422
Name: cholesterol , dtype: float64
```

```
In [139.. heartdata['cholesterol ']=heartdata['cholesterol '].fillna(value = heartdata.groupby(['sex'])['cholesterol '].tran:
```

Null values in the heart rate column are filled with the mode of heart rate, calculated within groups based on sex and chest pain.

```
In [139.. heartdata.groupby(['sex'])['heart rate'].mean()
```

```
Out[139.. sex
0    147.412791
1    143.455882
Name: heart rate, dtype: float64
```

```
In [140.. heartdata['heart rate']=heartdata['heart rate'].fillna(value = heartdata.groupby(['sex'])['heart rate'].tran:
```

Null values in the depression column are filled with the mode of depression, calculated within groups based on sex and chest pain.

```
In [140.. heartdata.groupby(['sex'])['depression '].mean()
```

```
Out[140.. sex
0    0.635673
1    0.858437
Name: depression , dtype: float64
```

```
In [140.. heartdata['depression ']=heartdata['depression '].fillna(value = heartdata.groupby(['sex'])['depression '].t
```

```
In [140.. pd.DataFrame({'Count':heartdata.isnull().sum()[heartdata.isnull().sum()>0], 'Percentage':heartdata.isnull().sum
```

```
Out[140.. Count Percentage
```

## Scaling

```
In [141.. from sklearn.preprocessing import MinMaxScaler
```

```
sc=MinMaxScaler()
```

```
heartdata_scaled=pd.DataFrame(sc.fit_transform(heartdata),columns=heartdata.columns,index=heartdata.index)
heartdata_scaled
```

```
Out[141..
```

	Age (age in year)	sex	chest pain	blood pressure	cholesterol	blood sugar	electrocardiographic	heart rate	exercise induced	depression	slope	ca	thal	c
0	0.714286	1.0	0.000000	0.679487	0.456181	1.0	1.0	0.584	0.0	0.613333	1.0	0.0	0.75	0.0
1	0.183673	1.0	0.666667	0.487179	0.518911	0.0	0.0	0.880	0.0	0.933333	1.0	0.0	0.00	0.0
2	0.265306	0.0	0.333333	0.487179	0.349170	0.0	1.0	0.760	0.0	0.373333	0.0	0.0	0.00	0.0
3	0.571429	1.0	0.333333	0.358974	0.467251	0.0	0.0	0.808	0.0	0.213333	0.0	0.0	0.00	0.0
4	0.591837	0.0	1.000000	0.358974	0.902675	0.0	0.0	0.688	1.0	0.160000	0.0	0.0	0.00	0.0
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
575	0.489796	1.0	1.000000	0.615385	0.577952	0.0	0.0	0.456	1.0	0.533333	0.5	0.0	1.00	1.0
576	0.306122	1.0	1.000000	0.615385	0.659133	0.0	0.0	0.464	1.0	0.533333	0.5	0.0	1.00	1.0
577	0.265306	1.0	1.000000	0.358974	0.836255	0.0	0.0	0.328	1.0	0.800000	0.5	0.0	1.00	1.0
578	0.326531	1.0	1.000000	0.551282	1.000000	0.0	0.0	0.464	0.0	0.000000	0.5	0.0	1.00	1.0
579	0.428571	1.0	1.000000	0.743590	0.415590	0.0	0.0	0.360	0.0	0.533333	0.5	0.0	1.00	1.0

580 rows × 14 columns

## Duplicate



```
In [141... #Dropped rows
df_dub = heartdata.drop_duplicates()

print(heartdata.shape, df_dub.shape)
#Dropped rows
```

(580, 14) (580, 14)

No duplicated value is detected.

## Data Splitting

```
In [141... x = heartdata_scaled.drop("c", axis=1)
y = heartdata_scaled.c # df["c"]

# x is my ind features
# y is my target
```

```
In [142... from sklearn.model_selection import train_test_split
# stratify = y
```

```
In [142... # from sklearn.model_selection import train_test_split

Xtrain, Xtest, Ytrain, Ytest = train_test_split(x, y, test_size = 0.2, random_state= 42, stratify=y)

print(Xtrain.shape)
print(Xtest.shape)
```

(464, 13)

(116, 13)

```
In [142... Ytest.value_counts()
```

```
Out[142... c
0.0    67
1.0    49
Name: count, dtype: int64
```

## Modeling

### Decision Tree

```
In [142... from sklearn.tree import DecisionTreeClassifier
```

#### Hyperparameter Analysis

```
In [143... hp_max=7

result=[]

for i in range(1, hp_max):

    DT=DecisionTreeClassifier(max_depth=i)
    DT.fit(Xtrain, Ytrain)
    result.append([i, DT.fit(Xtrain, Ytrain).score(Xtrain, Ytrain), accuracy_score(Ytest, DT.predict(Xtest))])
result
```

```
Out[143... [[1, 0.771551724137931, 0.75],
 [2, 0.7974137931034483, 0.8017241379310345],
 [3, 0.834051724137931, 0.7931034482758621],
 [4, 0.8556034482758621, 0.8189655172413793],
 [5, 0.8793103448275862, 0.8017241379310345],
 [6, 0.8987068965517241, 0.8017241379310345]]
```

Max depth= 4 was selected

```
In [143... DT = DecisionTreeClassifier(random_state = 42, max_depth = 4, min_samples_leaf = 1)

DT.fit(Xtrain, Ytrain)

pred = DT.predict(Xtest)
#Print the predicted values
pred
```

```
Out[143... array([0., 0., 1., 0., 0., 0., 1., 1., 0., 0., 0., 0., 0., 0., 1., 1., 0.,
        0., 0., 0., 1., 0., 0., 1., 0., 0., 1., 0., 1., 0., 0., 0., 0.,
        0., 0., 0., 0., 0., 1., 0., 0., 0., 0., 0., 0., 0., 0., 1., 0.,
        0., 0., 0., 0., 1., 0., 0., 1., 0., 0., 1., 0., 0., 0., 1., 0., 0.,
        1., 0., 0., 0., 0., 1., 1., 0., 0., 1., 1., 0., 1., 1., 0., 0., 1.,
        0., 0., 1., 0., 0., 0., 1., 1., 0., 0., 0., 1., 1., 1., 0., 0., 1.,
        0., 1., 0., 1., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.])
```

```
In [143... pred == Ytest
```

```
Out[143... 54      True
166      True
506      True
106      True
322      True
...
273      True
375     False
105      True
71       True
58       True
Name: c, Length: 116, dtype: bool
```

```
In [143... (pred == Ytest).mean()
# accuracy
```

```
Out[143... 0.8189655172413793
```

```
In [144... #Metrics:
acc = accuracy_score(Ytest, pred)
# Confusion matrix
# Confusion matrix
rec = recall_score(Ytest, pred)

pre = precision_score(Ytest, pred)

fm = f1_score(Ytest, pred)

conf = confusion_matrix(Ytest, pred)

print(acc, rec, pre, fm)
print(conf)

0.8189655172413793 0.6122448979591837 0.9375 0.7407407407407407
[[65  2]
 [19 30]]
```

```
In [144... print(classification_report(Ytest, pred))
```

	precision	recall	f1-score	support
0.0	0.77	0.97	0.86	67
1.0	0.94	0.61	0.74	49
accuracy			0.82	116
macro avg	0.86	0.79	0.80	116
weighted avg	0.84	0.82	0.81	116

```
In [144... tf= pd.DataFrame(data=result,columns=['Depth','Train','Test'])
tf
```

```
Out[144...
```

	Depth	Train	Test
0	1	0.771552	0.750000
1	2	0.797414	0.801724
2	3	0.834052	0.793103
3	4	0.855603	0.818966
4	5	0.879310	0.801724
5	6	0.898707	0.801724

```
In [144... tf.set_index('Depth',inplace=True)
tf
```

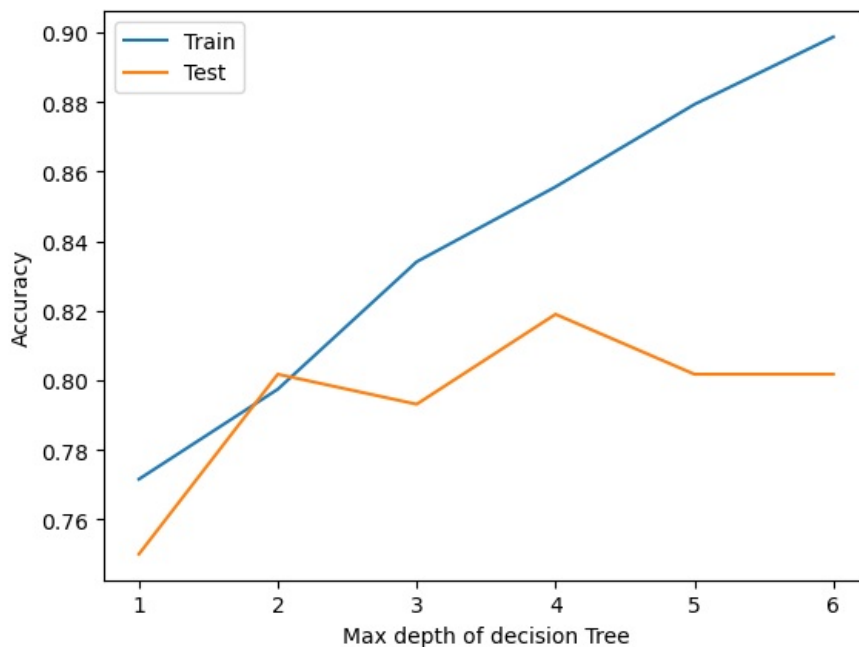
Out[144]:

## Train

## Test

Depth		
1	0.771552	0.750000
2	0.797414	0.801724
3	0.834052	0.793103
4	0.855603	0.818966
5	0.879310	0.801724
6	0.898707	0.801724

```
In [144]: tf.plot(kind='line',xlabel='Max depth of decision Tree',ylabel='Accuracy');
```



```
In [145... kf = KFold(10, shuffle = True, random_state = 42)
scorecross=cross_val_score(DT, x , y , cv = kf, scoring = 'accuracy' )
scorecross.mean()
```

Out[145]: 0.7948275862068965

## Concat Cross Val Score

```
In [145... pred = cross_val_predict(DT, x, y, cv = 10)
confusion matrix(y , pred)
```

```
Out[145]: array([[305, 32],
                  [ 88, 155]])
```

```
In [145]: # def metric_name(x,y):
from sklearn.model_selection import GridSearchCV
DT = DecisionTreeClassifier()

param = {'criterion':['gini', 'entropy'],
         'max_depth':[3,4,5,6,7],
         'min_samples_split':[3,4,5,6],
         'min_samples_leaf':[2,3,4]}

GS = GridSearchCV(DT, param ,cv = 10, scoring = "f1")
GS.fit(x, y)
# 交叉验证 网格搜索 参数 调优

# convex opt
# RandomSearch
# meta - Heuristics
```

```
Out[145... ▶ GridSearchCV ⓘ ?
  ▶ estimator: DecisionTreeClassifier
    ▶ DecisionTreeClassifier ?
```

```
In [145.. GS = GridSearchCV(DT, param ,cv = 10, scoring = "f1")
%time GS.fit(x, y)
```

CPU times: user 2.6 s, sys: 4.59 ms, total: 2.6 s  
Wall time: 2.62 s

```
Out[145.. ▶ GridSearchCV ⓘ ?
▶ estimator: DecisionTreeClassifier
    ▶ DecisionTreeClassifier ⓘ
```

```
In [145.. GS= GridSearchCV(DT, param ,cv = 10, scoring = "f1", n_jobs = -1)
%time GS.fit(x, y)
```

CPU times: user 278 ms, sys: 94.1 ms, total: 373 ms  
Wall time: 2.27 s

```
Out[145.. ▶ GridSearchCV ⓘ ?
▶ estimator: DecisionTreeClassifier
    ▶ DecisionTreeClassifier ⓘ
```

```
In [145.. GS.get_params()
```

```
Out[145.. {'cv': 10,
'error_score': nan,
'estimator__ccp_alpha': 0.0,
'estimator__class_weight': None,
'estimator__criterion': 'gini',
'estimator__max_depth': None,
'estimator__max_features': None,
'estimator__max_leaf_nodes': None,
'estimator__min_impurity_decrease': 0.0,
'estimator__min_samples_leaf': 1,
'estimator__min_samples_split': 2,
'estimator__min_weight_fraction_leaf': 0.0,
'estimator__monotonic_cst': None,
'estimator__random_state': None,
'estimator__splitter': 'best',
'estimator': DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',
max_depth=None, max_features=None, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_samples_leaf=1,
min_samples_split=2, min_weight_fraction_leaf=0.0,
monotonic_cst=None, random_state=None, splitter='best'),
'n_jobs': -1,
'param_grid': {'criterion': ['gini', 'entropy'],
'max_depth': [3, 4, 5, 6, 7],
'min_samples_split': [3, 4, 5, 6],
'min_samples_leaf': [2, 3, 4]},
'pre_dispatch': '2*n_jobs',
'refit': True,
'return_train_score': False,
'scoring': 'f1',
'verbose': 0}
```

```
In [145.. from sklearn import set_config
set_config(print_changed_only= False)
DT
```

```
Out[145.. ▼ DecisionTreeClassifier ⓘ ?
DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',
max_depth=None, max_features=None, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_samples_leaf=1,
min_samples_split=2, min_weight_fraction_leaf=0.0,
monotonic_cst=None, random_state=None, splitter='best')
```

```
In [146.. GS.best_params_
```

```
Out[146.. {'criterion': 'entropy',
'max_depth': 7,
'min_samples_leaf': 4,
'min_samples_split': 5}
```

```
In [146.. acc_DT=GS.best_score_
acc_DT
```

Out[146.. 0.7274418820616491

## KNN

```
In [146.. from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import cross_val_predict , cross_val_score
```

```
In [146.. KNN = KNeighborsClassifier(n_neighbors = 7, weights = "distance")
scores = cross_val_score(KNN, x , y , cv =10 , scoring = 'accuracy').mean()
```

```
In [146.. import matplotlib.pyplot as plt
from sklearn import metrics


from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.25,random_state=42)
```

```
In [146.. scores = []
k_range = list(range(1,20)) # >>[1,2,3,4,5, ... 25]
for i in k_range:
    knn = KNeighborsClassifier(n_neighbors=i)
    knn.fit(X_train, y_train)
    y_pred = knn.predict(X_test)
    scores.append(metrics.accuracy_score(y_test, y_pred))

plt.plot(k_range, scores)

plt.xlabel('Value of k for KNN')
plt.ylabel('Accuracy Score')
plt.title('Accuracy Scores for Values of k of k-Nearest-Neighbors')
```

Out[146.. Text(0.5, 1.0, 'Accuracy Scores for Values of k of k-Nearest-Neighbors')



```
In [146.. scores=np.array(scores)
scores
```

```
Out[146.. array([0.8       , 0.84137931, 0.85517241, 0.86896552, 0.85517241,
        0.87586207, 0.82758621, 0.84137931, 0.84827586, 0.84827586,
        0.83448276, 0.85517241, 0.84827586, 0.84827586, 0.84827586,
        0.86206897, 0.85517241, 0.84827586, 0.84137931])
```


```
In [146.. scores.mean()
```

Out[146.. 0.8475499092558985

```
In [147.. k_range=range(1,15)
weights_custom=["uniform","distance"]

param_grid = dict(n_neighbors=k_range, weights=weights_custom)
KNNGreed = GridSearchCV(KNN, param_grid ,cv = 5, scoring = "accuracy")
KNNGreed.fit(x, y)
```

Out[147..



In [147.. KNNGreed.cv\_results\_

```
Out[147.. {'mean_fit_time': array([0.00130301, 0.00113258, 0.00125022, 0.00170555, 0.00062089,
    0.00063262, 0.00067306, 0.00062976, 0.00083013, 0.00057101,
    0.00074201, 0.00055408, 0.00058331, 0.00056872, 0.00055904,
    0.00054975, 0.00062146, 0.00060825, 0.00072684, 0.00059495,
    0.00058346, 0.00073786, 0.00055947, 0.00056047, 0.00054674,
    0.00054183, 0.00054507, 0.00053878]),
  'std_fit_time': array([7.51845864e-04, 6.69308696e-04, 8.31153435e-04, 1.40086430e-03,
    3.64720894e-05, 5.87757845e-05, 5.76714781e-05, 5.00289675e-05,
    2.57128680e-04, 1.78702638e-05, 1.09200644e-04, 1.47673319e-05,
    3.17655873e-05, 1.16455701e-05, 6.22706244e-06, 7.76878809e-06,
    8.91563846e-05, 2.10378695e-05, 8.35451633e-05, 4.00601678e-05,
    3.21769347e-05, 1.07782751e-04, 7.80994636e-06, 2.84470035e-05,
    5.56369052e-06, 6.92908546e-06, 7.45667209e-06, 6.67980614e-06]),
  'mean_score_time': array([0.00474958, 0.0018559 , 0.00424805, 0.00187235, 0.00324321,
    0.00106449, 0.00278244, 0.00109725, 0.00334525, 0.00105429,
    0.00288463, 0.0010746 , 0.00264063, 0.00108118, 0.00260096,
    0.00111246, 0.00277367, 0.00118384, 0.00303378, 0.00120139,
    0.00283775, 0.00146914, 0.00267591, 0.00120034, 0.00265465,
    0.0012044 , 0.00267015, 0.0012167 ]),
  'std_score_time': array([2.04382325e-03, 8.95732229e-04, 2.87450404e-03, 1.01815412e-03,
    1.38869617e-03, 5.47967437e-05, 1.29049227e-04, 4.30970583e-05,
    7.86860086e-04, 3.09485316e-05, 1.86478387e-04, 3.09868583e-05,
    3.66841211e-05, 2.12443664e-05, 1.70900726e-05, 3.05369435e-05,
    1.79726092e-04, 3.72142610e-05, 2.22723516e-04, 6.02539931e-05,
    2.01336245e-04, 2.12744146e-04, 6.02486344e-05, 4.00232581e-05,
    2.47241791e-05, 3.05795785e-05, 2.29579141e-05, 1.55139979e-05]),
  'param_n_neighbors': masked_array(data=[1, 1, 2, 2, 3, 3, 4, 4, 5, 5, 6, 6, 7, 7, 8, 8, 9, 9,
    10, 10, 11, 11, 12, 12, 13, 13, 14, 14],
    mask=[False, False, False, False, False, False, False, False,
    False, False, False, False, False, False, False, False, False,
    False, False, False, False, False, False, False, False,
    False, False, False, False],
    fill_value='?',
    dtype=object),
  'param_weights': masked_array(data=['uniform', 'distance', 'uniform', 'distance',
    'uniform', 'distance', 'uniform', 'distance',
    'uniform', 'distance', 'uniform', 'distance',
    'uniform', 'distance', 'uniform', 'distance',
    'uniform', 'distance', 'uniform', 'distance',
    'uniform', 'distance', 'uniform', 'distance'],
    mask=[False, False, False, False, False, False, False, False,
    False, False, False, False, False, False, False, False,
    False, False, False, False, False, False, False, False,
    False, False, False, False],
    fill_value='?',
    dtype=object),
  'params': [{'n_neighbors': 1, 'weights': 'uniform'},
    {'n_neighbors': 1, 'weights': 'distance'},
    {'n_neighbors': 2, 'weights': 'uniform'},
    {'n_neighbors': 2, 'weights': 'distance'},
    {'n_neighbors': 3, 'weights': 'uniform'},
    {'n_neighbors': 3, 'weights': 'distance'},
    {'n_neighbors': 4, 'weights': 'uniform'},
    {'n_neighbors': 4, 'weights': 'distance'},
    {'n_neighbors': 5, 'weights': 'uniform'},
    {'n_neighbors': 5, 'weights': 'distance'},
    {'n_neighbors': 6, 'weights': 'uniform'},
    {'n_neighbors': 6, 'weights': 'distance'},
    {'n_neighbors': 7, 'weights': 'uniform'},
    {'n_neighbors': 7, 'weights': 'distance'},
    {'n_neighbors': 8, 'weights': 'uniform'},
    {'n_neighbors': 8, 'weights': 'distance'},
    {'n_neighbors': 9, 'weights': 'uniform'},
    {'n_neighbors': 9, 'weights': 'distance'},
    {'n_neighbors': 10, 'weights': 'uniform'},
    {'n_neighbors': 10, 'weights': 'distance'},
    {'n_neighbors': 11, 'weights': 'uniform'},
    {'n_neighbors': 11, 'weights': 'distance'},
    {'n_neighbors': 12, 'weights': 'uniform'},
    {'n_neighbors': 12, 'weights': 'distance'},
    {'n_neighbors': 13, 'weights': 'uniform'},
    {'n_neighbors': 13, 'weights': 'distance'}]
```

```
{'n_neighbors': 14, 'weights': 'uniform'},
{'n_neighbors': 14, 'weights': 'distance'}],
'split0_test_score': array([0.75          , 0.75          , 0.74137931, 0.75          , 0.72413793,
0.72413793, 0.71551724, 0.73275862, 0.76724138, 0.76724138,
0.74137931, 0.75          , 0.75          , 0.74137931, 0.75862069,
0.75          , 0.75862069, 0.75862069, 0.75862069, 0.75          ,
0.76724138, 0.76724138, 0.76724138, 0.75          , 0.75862069,
0.76724138, 0.75862069, 0.75862069]),
'split1_test_score': array([0.63793103, 0.63793103, 0.69827586, 0.63793103, 0.70689655,
0.70689655, 0.73275862, 0.69827586, 0.73275862, 0.73275862,
0.72413793, 0.71551724, 0.70689655, 0.70689655, 0.72413793,
0.73275862, 0.72413793, 0.72413793, 0.74137931, 0.73275862,
0.74137931, 0.74137931, 0.73275862, 0.73275862, 0.73275862,
0.73275862, 0.71551724, 0.73275862]),
'split2_test_score': array([0.79310345, 0.79310345, 0.77586207, 0.79310345, 0.79310345,
0.78448276, 0.79310345, 0.81896552, 0.81034483, 0.79310345,
0.79310345, 0.80172414, 0.80172414, 0.79310345, 0.80172414,
0.81034483, 0.81896552, 0.81896552, 0.80172414, 0.81034483,
0.81034483, 0.81034483, 0.81034483, 0.81034483, 0.81034483,
0.81034483, 0.81896552, 0.81034483]),
'split3_test_score': array([0.81896552, 0.81896552, 0.87068966, 0.81896552, 0.87068966,
0.86206897, 0.87931034, 0.87068966, 0.87931034, 0.87068966,
0.87931034, 0.86206897, 0.87068966, 0.86206897, 0.87931034,
0.86206897, 0.88793103, 0.88793103, 0.87931034, 0.87931034,
0.89655172, 0.89655172, 0.90517241, 0.87931034, 0.89655172,
0.88793103, 0.9137931 , 0.89655172]),
'split4_test_score': array([0.81896552, 0.81896552, 0.8362069 , 0.81896552, 0.82758621,
0.85344828, 0.85344828, 0.8362069 , 0.84482759, 0.84482759,
0.86206897, 0.85344828, 0.82758621, 0.8362069 , 0.8362069 ,
0.82758621, 0.81896552, 0.81896552, 0.84482759, 0.82758621,
0.84482759, 0.8362069 , 0.84482759, 0.82758621, 0.8362069 ,
0.82758621, 0.84482759, 0.8362069 ]),
'mean_test_score': array([0.7637931 , 0.7637931 , 0.78448276, 0.7637931 , 0.78448276,
0.7862069 , 0.79482759, 0.79137931, 0.80689655, 0.80172414,
0.8          , 0.79655172, 0.79137931, 0.78793103, 0.8          ,
0.79655172, 0.80172414, 0.80172414, 0.80517241, 0.8          ,
0.81206897, 0.81034483, 0.81206897, 0.8          , 0.80689655,
0.80517241, 0.81034483, 0.80689655]),
'std_test_score': array([0.06779174, 0.06779174, 0.06240331, 0.06779174, 0.06168463,
0.06390949, 0.06437295, 0.06506194, 0.05246595, 0.05026683,
0.06226024, 0.05707912, 0.05754593, 0.05759756, 0.05490236,
0.0483374 , 0.05639802, 0.05639802, 0.05160907, 0.05325326,
0.05517241, 0.05424873, 0.06007328, 0.05325326, 0.05785504,
0.05302951, 0.06874966, 0.05785504]),
'rank_test_score': array([26, 26, 24, 26, 25, 23, 19, 20, 5, 10, 13, 17, 20, 22, 13, 18, 10,
10, 8, 13, 1, 3, 1, 13, 5, 9, 3, 5], dtype=int32)}
```

In [147... KNNGreed.best\_params\_

Out[147... {'n\_neighbors': 11, 'weights': 'uniform'}

In [147... acc\_knn=KNNGreed.best\_score\_  
acc\_knn

Out[147... 0.8120689655172415

## Neural Network

In [147... from sklearn.neural\_network import MLPClassifier

In [147... MLP = MLPClassifier()

```
scores = cross_val_score(MLP, x , y , cv =10 , scoring = 'accuracy')
scores.mean()
```

Out[147... 0.8344827586206897

In [147... from sklearn.model\_selection import GridSearchCV

```
MLP = MLPClassifier(random_state = 42)
```

```
param = {"activation" : ["relu" , "logistic" , "tanh"],
         "hidden_layer_sizes":[(10), (20), (20,30)],
         "max_iter" : [10, 50, 100, 200],
         # "solver": ["sgd", "adam"],
         "learning_rate_init": [0.01, 0.001, 0.025]}
```

```
GS = GridSearchCV(MLP, param, cv = 10)
GS.fit(x , y)
```

```
Out[147... GridSearchCV
  estimator: MLPClassifier
    MLPClassifier
```

```
In [147... GS.best_params_
```

```
Out[147... {'activation': 'relu',
  'hidden_layer_sizes': (20, 30),
  'learning_rate_init': 0.025,
  'max_iter': 10}
```

```
In [147... acc_nn=GS.best_score_
acc_nn
```

```
Out[147... 0.85
```

## Logistic Regression

```
In [148... from sklearn.linear_model import LogisticRegression
```

```
In [148... import statsmodels.api as sm
```

```
In [148... logreg = LogisticRegression( random_state = 42).fit(X_train,y_train)
logreg

print("Training set score: {:.3f}".format(logreg.score(X_train,y_train)))
print("Test set score: {:.3f}".format(logreg.score(X_test,y_test)))

import statsmodels.api as sm
# x = sm.add_constant(x)

logit_model=sm.Logit(y,x)
result=logit_model.fit()
print(result.summary())
```

```
Training set score: 0.832
Test set score: 0.862
Optimization terminated successfully.
    Current function value: 0.392219
    Iterations 7
```

Logit Regression Results						
=====						
Dep. Variable:	c	No. Observations:	580			
Model:	Logit	Df Residuals:	567			
Method:	MLE	Df Model:	12			
Date:	Wed, 16 Oct 2024	Pseudo R-squ.:	0.4232			
Time:	15:06:41	Log-Likelihood:	-227.49			
converged:	True	LL-Null:	-394.37			
Covariance Type:	nonrobust	LLR p-value:	3.696e-64			
=====						
	coef	std err	z	P> z	[0.025	0.975]
-----						
Age (age in year)	-2.3886	0.667	-3.580	0.000	-3.696	-1.081
sex	0.0608	0.356	0.171	0.864	-0.636	0.758
chest pain	1.1362	0.370	3.069	0.002	0.411	1.862
blood pressure	-0.2886	0.569	-0.507	0.612	-1.404	0.826
cholesterol	0.0875	0.605	0.145	0.885	-1.098	1.273
blood sugar	0.2550	0.389	0.656	0.512	-0.507	1.017
electrocardiographic	0.5280	0.303	1.745	0.081	-0.065	1.121
heart rate	-4.0220	0.566	-7.111	0.000	-5.131	-2.913
exercise induced	1.0964	0.284	3.860	0.000	0.540	1.653
depression	2.9783	0.567	5.252	0.000	1.867	4.090
slope	-0.3913	0.509	-0.769	0.442	-1.389	0.606
ca	3.6927	0.695	5.310	0.000	2.330	5.056
thal	1.4656	0.355	4.132	0.000	0.770	2.161
=====						

```
In [148... acc_lr=0.82
```

## Support Vector Machines

```
In [148... from sklearn.svm import SVC
```



```
In [149.. SVM = SVC()
score = cross_val_score(SVM , x, y, cv = 10 )
score.mean()
```

Out[149.. 0.8396551724137931

```
In [150.. SVM = SVC(random_state = 42 , class_weight = {0:0.4 , 1:0.6})

param = [{"kernel" : ["linear"] , "C" : [0.01 , 0.1, 1, 10, 100]},
        {"kernel" : ["rbf"], "gamma" : [0.01, 0.1, 0.2, 0.3], "C": [0.01 , 0.1, 1, 10, 100]},
        {"kernel" : ["poly"], "degree": [2], "C": [0.01, 0.1, 1, 10, 100]}]

##### HINT
#이름을 바꿔서 GridSearchCV로 바꿔서 실행
#이름을 바꿔서 GridSearchCV로 바꿔서 실행

GS = GridSearchCV(SVM, param, cv = 5, scoring = "accuracy" )

GS.fit(x , y)
```

Out[150..

▶ **GridSearchCV** ⓘ ?

▶ **estimator: SVC**

▶ **SVC** ?

```
In [150.. GS.best_params_
```

Out[150.. {'C': 1, 'kernel': 'linear'}

```
In [150.. acc_svm=GS.best_score_
acc_svm
```

Out[150.. 0.8172413793103448

# Naive Bayes

```
In [150.. from sklearn.naive_bayes import GaussianNB , MultinomialNB
```

```
In [151.. GNB = GaussianNB()
scores = cross_val_score(GNB, x , y , cv =10 , scoring = 'accuracy')
acc_nbg=scores.mean()
acc_nbg
```

Out[151.. 0.8206896551724139

```
In [151.. #이름을 바꿔서 GridSearchCV로 바꿔서 실행
MNB = MultinomialNB()
scores = cross_val_score(MNB, x , y , cv =10 , scoring = 'accuracy')
acc_nbm=scores.mean()
acc_nbm
```

Out[151.. 0.789655172413793

```
In [151.. pd.DataFrame(
    [acc_DT, acc_nn, acc_knn, acc_lr, acc_nbg, acc_nbm, acc_svm],
    index=['Decision Tree', 'Neural Network', 'K-Nearest Neighbor', 'Logistic Regression', 'Naive Bayes Gaussian', 'Naive Bayes Multi', 'Support Vector Machine'],
    columns=['Cross Val. Accuracy'])
```

Out[151..

	Cross Val. Accuracy
Decision Tree	0.727442
Neural Network	0.850000
K-Nearest Neighbor	0.812069
Logistic Regression	0.820000
Naive Bayes Gaussian	0.820690
Naive Bayes Multi	0.789655
Support Vector Machine	0.817241